

# COMP6771

# Advanced C++ Programming

Week 5.2

Smart Pointers

# In this lecture

## Why?

- Managing unnamed / heap memory can be dangerous, as there is always the chance that the resource is not released / free'd properly. We need solutions to help with this.

## What?

- Smart pointers
- Unique pointer, shared pointer
- Partial construction

# Recap: RAI - Making unnamed objects safe

Don't use the new / delete keyword in your own code

We are showing for demonstration purposes

```
1 // myintpointer.h
2
3 class MyIntPtr {
4 public:
5     // This is the constructor
6     MyIntPtr(int* value);
7
8     // This is the destructor
9     ~MyIntPtr();
10
11 int* GetValue();
12
13 private:
14     int* value_;
15 };
```

```
1 // myintpointer.cpp
2 #include "myintpointer.h"
3
4 MyIntPtr::MyIntPtr(int* value): value_{value} {}
5
6 int* MyIntPtr::GetValue() {
7     return value_
8 }
9
10 MyIntPtr::~MyIntPtr() {
11     // Similar to C's free function.
12     delete value_;
13 }
```

```
1 void fn() {
2     // Similar to C's malloc
3     MyIntPtr p{new int{5}};
4     // Copy the pointer;
5     MyIntPtr q{p.GetValue()};
6     // p and q are both now destructed.
7     // What happens?
8 }
```

demo551-safepointer.cpp

# Smart Pointers

- Ways of wrapping unnamed (i.e. raw pointer) heap objects in named stack objects so that object lifetimes can be managed much easier
- Introduced in C++11
- Usually two ways of approaching problems:
  - `unique_ptr` + raw pointers ("observers")
  - `shared_ptr` + `weak_ptr`/raw pointers

*doesn't have copy constructor (deleted from the type)*

Type	Shared ownership	Take ownership
<u><code>std::unique_ptr&lt;T&gt;</code></u>	No	Yes
<i>int* i;</i> raw pointers	No	No
<code>std::shared_ptr&lt;T&gt;</code>	Yes	Yes
<code>std::weak_ptr&lt;T&gt;</code>	No	No

# Unique pointer


```
2 #include <memory>
3
4 auto main() -> int {
5     auto my_heap_object = new std::string{"Hi"}; // HEAP
6     {
7         You, seconds ago • Uncommitted changes
8         auto up1 = std::unique_ptr<std::string>(my_heap_object);
9         std::cout << "What is this???" << (*up1).size() << "\n";
10    }
11    std::cout << "What is this???" << (*my_heap_object).size() << "\n";
12 }
```

- **std::unique\_ptr<T>**

- The unique pointer owns the object
- When the unique pointer is destructed, the underlying object is too

- **raw pointer (observer)**

- Unique Ptr may have many observers
- This is an appropriate use of raw pointers (or references) in C++
- Once the original pointer is destructed, you must ensure you don't access the raw pointers (no checks exist)
- These observers **do not** have ownership of the pointer



Also note the use of 'nullptr' in C++ instead of NULL

Why no copy constructor / assignment:

# Unique pointer: Usage

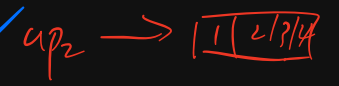

(case 2:

case 1:




Doesn't make sense, one go out of scope, another point to freed heap memory.

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     auto up1 = std::unique_ptr<int>{new int};
6     auto up2 = up1; // no copy constructor
7     std::unique_ptr<int> up3;
8     up3 = up2; // no copy assignment
9
10    { up3.reset(up1.release()); // OK
11      auto up4 = std::move(up3); // OK
12      std::cout << up4.get() << "\n";
13      std::cout << *up4 << "\n";
14      std::cout << *up1 << "\n";
15    }
```



Waste of memory.

Move:



reset: replace manage object (move assignment)

get(): return the actual pointer

demo552-unique1.cpp

# Observer Ptr: Usage

```
1 #include <memory>
2 #include <iostream>
3
4 int main() {
5     auto up1 = std::unique_ptr<int>{new int{0}};
6     *up1 = 5;
7     std::cout << *up1 << "\n";
8     auto op1 = up1.get();
9     *op1 = 6;
10    std::cout << *op1 << "\n";
11    up1.reset();
12    std::cout << *op1 << "\n";
13 }
```

demo553-observer.cpp

Can we remove "new"  
completely?

Yes.

# Unique Ptr Operators

This method avoids the need for "new". It has other benefits that we will explore.

```
1 #include <iostream>
2 #include <memory>
3
4 auto main() -> int {
5     // 1 - Worst - you can accidentally own the resource multiple
6     // times, or easily forget to own it.
7     // auto* silly_string = new std::string{"Hi"};
8     // auto up1 = std::unique_ptr<std::string>(silly_string);
9     // auto up11 = std::unique_ptr<std::string>(silly_string);
10
11     // 2 - Not good - requires actual thinking about whether there's a leak.
12     auto up2 = std::unique_ptr<std::string>(new std::string("Hello"));
13
14     // 3 - Good - no thinking required.
15     auto up3 = std::make_unique<std::string>("Hello");
16
17     std::cout << *up2 << "\n";
18     std::cout << *up3 << "\n";
19     // std::cout << *(up3.get()) << "\n";
20     // std::cout << up3->size();
21 }
```


← SAME thing

demo554-unique2.cpp

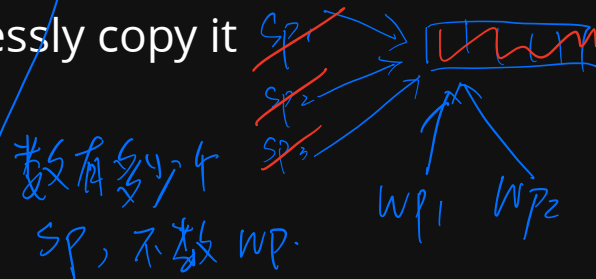
- <https://stackoverflow.com/questions/37514509/advantages-of-using-stdmake-unique-over-new-operator>
- <https://stackoverflow.com/questions/20895648/difference-in-make-shared-and-normal-shared-Ptr-in-c>



# Shared pointer



- `std::shared_ptr<T>`
- Several shared pointers share ownership of the object
  - A reference counted pointer
  - When a shared pointer is destructed, if it is the only shared pointer left pointing at the object, then the **object is destroyed**
  - May also have many observers
    - Just because the pointer has shared ownership doesn't mean the observers should get ownership too - don't mindlessly copy it
- `std::weak_ptr<T>`
  - Weak pointers are used with share pointers when:
    - You don't want to add to the reference count
    - You want to be able to check if the underlying data is still valid before using it.



# Shared pointer: Usage

Stack  
X →  
y →  
Heap  
[5]

```
1 #include <iostream>
2 #include <memory>
3
4 auto main() -> int {
5     auto x = std::make_shared<int>(5);
6     std::cout << "use count: " << x.use_count() << "\n";
7     std::cout << "value: " << *x << "\n";
8     x.reset(); // Memory still exists, due to y.
9     std::cout << "use count: " << y.use_count() << "\n";
10    std::cout << "value: " << *y << "\n";
11    y.reset(); // Deletes the memory, since
12    // no one else owns the memory
13    std::cout << "use count: " << x.use_count() << "\n";
14    std::cout << "value: " << *y << "\n";
15 }
```

auto y = x;

demo555-shared.cpp

undefined behavior (dereferenced freed heap)

Can we remove "new" completely?

# Weak Pointer: Usage

```
1 #include <iostream>
2 #include <memory>
3
4 auto main() -> int {
5     auto x = std::make_shared<int>(1);
6
7     auto wp = std::weak_ptr<int>(x); // x owns the memory
8
9     auto y = wp.lock();
10    if (y != nullptr) { // x and y own the memory ✓
11        // Do something with y
12        std::cout << "Attempt 1: " << *y << '\n';
13    }
14 }
```

*(lock) return  
a shared pointer  
if object  
exist, or nullptr.*

demo556-weak.cpp

# △ When to use which type

- **Unique pointer vs shared pointer**

- You almost always want a unique pointer over a shared pointer
- Use a shared pointer if either:
  - An object has multiple owners, **and you don't know which one will stay around the longest**
  - You need temporary ownership (outside scope of this course)
  - This is very rare

# Smart pointer examples

- Linked list
- Doubly linked list
- Tree
- DAG (mutable and non-mutable)
- Graph (mutable and non-mutable)
- Twitter feed with multiple sections (eg. my posts, popular posts)

## “Leak freedom in C++” poster

Strategy	Natural examples	Cost	Rough frequency
<b>1. Prefer scoped lifetime</b> by default (locals, members)	Local and member objects – directly owned	Zero: Tied directly to another lifetime	<div> <div>O(80%) of objects</div> <div>O(20%) of objects</div> </div>
<b>2. Else prefer <code>make_unique</code> &amp; <code>unique_ptr</code> or a container</b> , if the object must have its own lifetime (i.e., heap) and ownership can be unique w/o owning cycles	Implementations of trees, lists	Same as new/delete & malloc/free <b>Automates</b> simple heap use in a library	
<b>3. Else prefer <code>make_shared</code> &amp; <code>shared_ptr</code></b> , if the object must have its own lifetime (i.e., heap) and shared ownership w/o owning cycles	Node-based DAGs, incl. trees that share out references	Same as manual reference counting (RC) <b>Automates</b> shared object use in a library	

**Don't use owning raw `*`'s == don't use explicit `delete`**

**Don't create ownership cycles** across modules by owning “upward” (violates layering)

Use `weak_ptr` to break cycles

# Stack unwinding

- Stack unwinding is the process of exiting the stack frames until we find an exception handler for the function
- This calls any destructors on the way out
  - Any resources not managed by destructors won't get freed up
  - If an exception is thrown during stack unwinding, `std::terminate` is called

Not safe

```
1 void g() {
2     throw std::runtime_error("");
3 }
4
5 int main() {
6     auto ptr = new int{5};
7     g();
8     // Never executed.
9     delete ptr;
10 }
```

Not safe

```
1 void g() {
2     throw std::runtime_error("");
3 }
4
5 int main() {
6     auto ptr = new int{5};
7     g();
8     auto uni = std::unique_ptr<int>(ptr);
9 }
```

Safe

```
1 void g() {
2     throw std::runtime_error("");
3 }
4
5 int main() {
6     auto ptr = std::make_unique<int>(5);
7     g();
8 }
```

# Exceptions & Destructors

- During stack unwinding, `std::terminate()` will be called if an exception leaves a destructor
- The resources may not be released properly if an exception leaves a destructor
- All exceptions that occur inside a destructor should be handled inside the destructor
- Destructors usually don't throw, and need to explicitly opt in to throwing
  - STL types don't do that



# Partial construction

Spot the bug

- What happens if an exception is thrown halfway through a constructor?
  - The C++ standard: "An object that is partially constructed or partially destroyed will have destructors executed for all of its fully constructed subobjects"
  - A destructor is not called for an object that was partially constructed
  - Except for an exception thrown in a constructor that delegates (why?)

```
1 #include <exception>
2
3 class my_int {
4 public:
5     my_int(int const i) : i_{i} {
6         if (i == 2) {
7             throw std::exception();
8         }
9     }
10 private:
11     int i_;
12 };
13
14 class unsafe_class {
15 public:
16     unsafe_class(int a, int b)
17         : a_{new my_int{a}}
18         , b_{new my_int{b}}
19     {}
20
21 ~unsafe_class() {
22     delete a_;
23     delete b_;
24 }
25 private:
26     my_int* a_;
27     my_int* b_;
28 };
29
30 int main() {
31     auto a = unsafe_class(1, 2);
32 }
```

demo557-bad.cpp

Stack: a deleted, but object pointed  
never been freed (memory leak!!!)  
V ← not constructed  
partially constructed.

# Partial construction: Solution

- Option 1: Try / catch in the constructor
  - Very messy, but works (if you get it right...)
  - Doesn't work with initialiser lists (needs to be in the body)
- Option 2:
  - An object managing a resource should initialise the resource last
    - The resource is only initialised when the whole object is
    - Consequence: An object can only manage one resource
    - If you want to manage multiple resources, instead manage several wrappers , which each manage one resource

```
1 #include <exception>
2 #include <memory>
3
4 class my_int {
5 public:
6     my_int(int const i)
7         : i_{i} {
8         if (i == 2) {
9             throw std::exception();
10        }
11    }
12 private:
13     int i_;
14 };
15
16 class safe_class {
17 public:
18     safe_class(int a, int b)
19         : a_(std::make_unique<my_int>(a))
20         , b_(std::make_unique<my_int>(b))
21     {}
22 private:
23     std::unique_ptr<my_int> a_;
24     std::unique_ptr<my_int> b_;
25 };
26
27 int main() {
28     auto a = safe_class(1, 2);
29 }
```

demo558-partial1.cpp

# Feedback

